

Kernel Methods (and Some Neural Networks) for Rough PDEs

Ricardo Baptista¹ Edoardo Calvello¹ **Matthieu Darcy**¹ Houman Owhadi¹
Andrew M. Stuart¹ Xianjin Yang¹

¹Department of Computing and Mathematical Sciences
California Institute of Technology

Workshop on digital twins for inverse problems in Earth science
CIRM 2024



The problem

We want to **solve with machine learning rough PDEs** of the form

$$\begin{aligned}\mathcal{P}(u) &= \xi, & x \in \Omega, \\ u &= 0, & x \in \partial\Omega\end{aligned}\tag{RPDE}$$

- $\mathcal{P} : H_0^t(\Omega) \rightarrow H^{-s}$ is (non-linear) differential operator.

Canonical example: $\mathcal{P}(u) = -\Delta u + f(u)$.

- $u^* \in H_0^t(\Omega)$ is the **solution** and $\xi \in H^{-s}(\Omega)$ is the **forcing term**.

Talk summary

The problem

We want to **solve with machine learning rough PDEs** of the form

$$\begin{aligned}\mathcal{P}(u) &= \xi, & x \in \Omega, \\ u &= 0, & x \in \partial\Omega\end{aligned}\tag{RPDE}$$

- $\mathcal{P} : H_0^t(\Omega) \rightarrow H^{-s}$ is (non-linear) differential operator.

Canonical example: $\mathcal{P}(u) = -\Delta u + f(u)$.

- $u^* \in H_0^t(\Omega)$ is the **solution** and $\xi \in H^{-s}(\Omega)$ is the **forcing term**.

Hierarchy of spaces

$$\underbrace{H_0^t(\Omega)}_{\text{Functions with } t \text{ derivatives}} \subset L^2(\Omega) \subset \underbrace{H^{-s}}_{\text{Dual space of } H_0^s}$$

Talk summary

The problem

We want to **solve with machine learning rough PDEs** of the form

$$\begin{aligned}\mathcal{P}(u) &= \xi, & x \in \Omega, \\ u &= 0, & x \in \partial\Omega\end{aligned}\tag{RPDE}$$

- $\mathcal{P} : H_0^t(\Omega) \rightarrow H^{-s}$ is (non-linear) differential operator.

Canonical example: $\mathcal{P}(u) = -\Delta u + f(u)$.

- $u^* \in H_0^t(\Omega)$ is the **solution** and $\xi \in H^{-s}(\Omega)$ is the **forcing term**.

The difficulty

Solving (RPDE) is difficult because the forcing term is rough $\xi \notin L^2$ and, as a result, the solution u^* is also irregular/rough.

Talk summary

The problem

We want to **solve with machine learning rough PDEs** of the form

$$\begin{aligned}\mathcal{P}(u) &= \xi, & x \in \Omega, \\ u &= 0, & x \in \partial\Omega\end{aligned}\tag{RPDE}$$

- $\mathcal{P} : H_0^t(\Omega) \rightarrow H^{-s}$ is (non-linear) differential operator.

Canonical example: $\mathcal{P}(u) = -\Delta u + f(u)$.

- $u^* \in H_0^t(\Omega)$ is the **solution** and $\xi \in H^{-s}(\Omega)$ is the **forcing term**.

The solution

We propose a kernel based method for solving rough PDEs using **negative Sobolev norms** and **weak measurements**, with **provable convergence**.

Outline

- ① Introduction and motivation
- ② Machine learning for smooth PDEs
- ③ Machine learning for rough PDEs
- ④ Convergence results
- ⑤ Numerical results

Table of Contents

- 1 Introduction and motivation
- 2 Machine learning for smooth PDEs
- 3 Machine learning for rough PDEs
- 4 Convergence results
- 5 Numerical results

Motivation: Stochastic Partial Differential Equations

Stochastic Partial Differential Equations (SPDEs) are PDEs that include **random fluctuations**. A general class of interest are semi-linear SPDEs with additive noise:

$$\partial_t u(t, \mathbf{x}) = \Delta u(\mathbf{x}, t) + f(u(t, \mathbf{x})) + \xi(t, \mathbf{x}) \quad (\text{SL-SPDE})$$

where:

ξ is **stochastic forcing term** ex: space-time white noise.

Motivation: Stochastic Partial Differential Equations

Stochastic Partial Differential Equations (SPDEs) are PDEs that include **random fluctuations**. A general class of interest are semi-linear SPDEs with additive noise:

$$\partial_t u(t, \mathbf{x}) = \Delta u(\mathbf{x}, t) + f(u(t, \mathbf{x})) + \xi(t, \mathbf{x}) \quad (\text{SL-SPDE})$$

where:

ξ is **stochastic forcing term** ex: space-time white noise.

Examples

- Phase field models (Allen-Cahn equation).
- Mathematical biology (Nagumo equation).
- Filtering and sampling (Kushner–Stratonovich).

Motivation: Solving SPDEs

Numerically solving an SPDE of the form

$$\partial_t u(t, \mathbf{x}) = \Delta u(\mathbf{x}, t) + f(u(t, \mathbf{x})) + \xi(t, \mathbf{x}) \quad (\text{SL-SPDE})$$

typically involves drawing samples of ξ and solving (SL-SPDE) for that realization.

Motivation: Solving SPDEs

Numerically solving an SPDE of the form

$$\partial_t u(t, \mathbf{x}) = \Delta u(\mathbf{x}, t) + f(u(t, \mathbf{x})) + \xi(t, \mathbf{x}) \quad (\text{SL-SPDE})$$

typically involves drawing samples of ξ and solving (SL-SPDE) for that realization.

Difficulty

Solving (SL-SPDE) is difficult because

- 1 The forcing term $\xi \notin L^2$ a.s. and is not pointwise defined.
- 2 The solution u^* is irregular/rough.

This motivates us to develop methods to solve PDEs with very rough forcing terms and/or irregular solutions.

Table of Contents

- 1 Introduction and motivation
- 2 Machine learning for smooth PDEs**
- 3 Machine learning for rough PDEs
- 4 Convergence results
- 5 Numerical results

Solving PDEs with ML: the usual way

Machine learning methods solve PDEs by **minimizing a physics informed loss**:

$$u^\dagger := \arg \min_{u \in \mathcal{S}} \underbrace{\|\mathcal{P}(u) - \xi\|_{L^2(\Omega)}^2}_{\text{PDE data}} + \underbrace{\|u\|_{L^2(\partial\Omega)}^2}_{\text{Boundary data}} + \underbrace{\gamma \mathcal{R}(u)}_{\text{Regularization}} \quad \text{Infinite data}$$

Solving PDEs with ML: the usual way

Machine learning methods solve PDEs by **minimizing a physics informed loss**:

$$u^\dagger := \arg \min_{u \in \mathcal{S}} \underbrace{\|\mathcal{P}(u) - \xi\|_{L^2(\Omega)}^2}_{\text{PDE data}} + \underbrace{\|u\|_{L^2(\partial\Omega)}^2}_{\text{Boundary data}} + \underbrace{\gamma \mathcal{R}(u)}_{\text{Regularization}} \quad \text{Infinite data}$$

⇓ Discretization

$$\approx \arg \min_{u \in \mathcal{S}} \underbrace{\frac{1}{N_\Omega} \sum_{i=1}^{N_\Omega} |\mathcal{P}(u)(x_i) - \xi(x_i)|^2}_{\text{PDE data approximation}} + \underbrace{\frac{1}{N_{\partial\Omega}} \sum_{j=1}^{N_{\partial\Omega}} |u(x_j)|^2}_{\text{Boundary approximation}} + \gamma \mathcal{R}(u) \quad \text{Finite data}$$

Three difficulties with Rough PDEs

- 1 The L^2 norm is not appropriate.
- 2 Solves the PDE pointwise.
- 3 Requires that u^* may be well approximated by the class \mathcal{S} .

Solving PDEs with ML: the usual way

Machine learning methods solve PDEs by **minimizing a physics informed loss**:

$$u^\dagger := \arg \min_{u \in \mathcal{S}} \underbrace{\|\mathcal{P}(u) - \xi\|_{L^2(\Omega)}^2}_{\text{PDE data}} + \underbrace{\|u\|_{L^2(\partial\Omega)}^2}_{\text{Boundary data}} + \underbrace{\gamma \mathcal{R}(u)}_{\text{Regularization}} \quad \text{Infinite data}$$

⇓ Discretization

$$\approx \arg \min_{u \in \mathcal{S}} \underbrace{\frac{1}{N_\Omega} \sum_{i=1}^{N_\Omega} |\mathcal{P}(u)(x_i) - \xi(x_i)|^2}_{\text{PDE data approximation}} + \underbrace{\frac{1}{N_{\partial\Omega}} \sum_{j=1}^{N_{\partial\Omega}} |u(x_j)|^2}_{\text{Boundary approximation}} + \gamma \mathcal{R}(u) \quad \text{Finite data}$$

Two main approaches to solving PDEs with ML

- 1 Physics Informed Neural networks [Raissi et al.; Tancik et al.]: the class \mathcal{S} is a parametric class of deep neural networks u_θ .
- 2 Kernel methods/Gaussian Processes [Chen et al.]: the class \mathcal{S} is a *Reproducing Kernel Hilbert Space* (RKHS).

Solving PDEs: the RKHS approach

RKHS approach [Chen et al.]: select u^\dagger in a RKHS

Solve the non-linear least-squares problem:

$$u^\dagger = \arg \min_{u \in \mathcal{H}_K} \frac{1}{N_\Omega} \sum_{i=1}^{N_\Omega} |\mathcal{P}(u)(x_i) - \xi(x_i)|^2 + \frac{1}{N_{\partial\Omega}} \sum_{j=1}^{N_{\partial\Omega}} |u(x_j)|^2 + \underbrace{\gamma \|u\|_{\mathcal{H}_K}^2}_{\text{RKHS regularization}}$$

Interpreted as a **MAP estimator** of a Gaussian Process conditioned on non-linear measurements.

Solving PDEs: the RKHS approach

RKHS approach [Chen et al.]: select u^\dagger in a RKHS

Solve the non-linear least-squares problem:

$$u^\dagger = \arg \min_{u \in \mathcal{H}_K} \frac{1}{N_\Omega} \sum_{i=1}^{N_\Omega} |\mathcal{P}(u)(x_i) - \xi(x_i)|^2 + \frac{1}{N_{\partial\Omega}} \sum_{j=1}^{N_{\partial\Omega}} |u(x_j)|^2 + \underbrace{\gamma \|u\|_{\mathcal{H}_K}^2}_{\text{RKHS regularization}}$$

Interpreted as a **MAP estimator** of a Gaussian Process conditioned on non-linear measurements.

The kernel approach has strong theoretical guarantees:

- Provable convergence (**under the condition that** $u^* \in \mathcal{H}_K$).
- Error estimates [Batlle et al.].
- Bayesian interpretation provides uncertainty quantification.


Table of Contents

- 1 Introduction and motivation
- 2 Machine learning for smooth PDEs
- 3 Machine learning for rough PDEs**
- 4 Convergence results
- 5 Numerical results

Challenges and solutions

Three difficulties with Rough PDEs

- 1 The L^2 norm is not appropriate.
- 2 Solves the PDE pointwise.
- 3 Requires $u^* \in \mathcal{H}_K$ for the convergence theory.



Solutions

Our main contributions are to solve the three main difficulties:

- 1 Using a **negative Sobolev norm** H^{-s} instead of the usual L^2 norm.
- 2 **Efficient approximation** of the H^{-s} norm with weak measurements.
- 3 **Convergence results** of the method without $u^* \in \mathcal{H}_k$.

A Novel Approach to Solving Rough PDEs

Modified loss adapted to the roughness of the forcing term:

$$u^\gamma = \arg \min_{u \in \mathcal{H}_K} \|\mathcal{P}(u) - \xi\|_{H^{-s}}^2 + \|u\|_{L^2(\partial\Omega)}^2 + \gamma \|u\|_{\mathcal{H}_K}^2 \quad \text{Infinite data}$$

⇓ Discretization

$$u^{\gamma, N, M} = \arg \min_{u \in \mathcal{H}_K} \|\mathcal{P}(u) - \xi\|_{\Phi^N}^2 + \sum_{i=1}^m |u(x_i)|^2 + \gamma \|u\|_{\mathcal{H}_K}^2 \quad \text{Finite data}$$

We provide a computationally efficient way to discretize the negative Sobolev norms through a test space $\Phi^N := \text{span}\{\varphi_i\}_{i=1}^N$ (ex: Fourier, finite element, Haar ...).

¹See also [wPINN]

A Novel Approach to Solving Rough PDEs

Modified loss adapted to the roughness of the forcing term:

$$u^\gamma = \arg \min_{u \in \mathcal{H}_K} \|\mathcal{P}(u) - \xi\|_{H^{-s}}^2 + \|u\|_{L^2(\partial\Omega)}^2 + \gamma \|u\|_{\mathcal{H}_K}^2 \quad \text{Infinite data}$$

⇓ Discretization

$$u^{\gamma, N, M} = \arg \min_{u \in \mathcal{H}_K} \|\mathcal{P}(u) - \xi\|_{\Phi^N}^2 + \sum_{i=1}^m |u(x_i)|^2 + \gamma \|u\|_{\mathcal{H}_K}^2 \quad \text{Finite data}$$

We provide a computationally efficient way to discretize the negative Sobolev norms through a test space $\Phi^N := \text{span}\{\varphi_i\}_{i=1}^N$ (ex: Fourier, finite element, Haar ...).

This method can be interpreted as solving the PDE in weak form¹:

$$[\mathcal{P}(u), \varphi_i] = [\xi, \varphi_i] \quad i = 1, \dots, N.$$

¹See also [wPINN]

Solving rough PDEs with kernel methods

New objective function:

$$u^{\gamma, N, M} = \arg \min_{u \in \mathcal{H}_K} \|\mathcal{P}(u) - \xi\|_{\Phi^N}^2 + \sum_{i=1}^m |u(x_i)|^2 + \gamma \|u\|_{\mathcal{H}_K}^2$$

- Minimized through a Gauss-Newton formulated on function space (very fast, converges in < 10 steps).
- Closed form solution for linear problems.

Solving rough PDEs with kernel methods

New objective function:

$$u^{\gamma, N, M} = \arg \min_{u \in \mathcal{H}_K} \|\mathcal{P}(u) - \xi\|_{\Phi^N}^2 + \sum_{i=1}^m |u(x_i)|^2 + \gamma \|u\|_{\mathcal{H}_K}^2$$

- Minimized through a Gauss-Newton formulated on function space (very fast, converges in < 10 steps).
- Closed form solution for linear problems.

We can also minimize this loss with a PINN u_θ :

- Minimized through gradient descent (much slower than Gauss-Newton).
- Perform poorly on linear problems but are good on non-linear problems.
- Typically uses no regularization ($\gamma = 0$).
- Requires a random Fourier layer [Tancik et al.] to learn the high frequencies.

Table of Contents

- 1 Introduction and motivation
- 2 Machine learning for smooth PDEs
- 3 Machine learning for rough PDEs
- 4 Convergence results**
- 5 Numerical results

Convergence: Assumptions

Assumptions on the PDE

- The operator $\mathcal{P} : H_0^t \rightarrow H^{-s}$ is continuous.
- The solution operator is (locally) stable

$$\|u - u^*\|_{H_0^t} \leq C \|\mathcal{P}(u) - \xi\|_{H^{-s}}$$

Assumptions on the method

- The space Φ^N is dense in H^{-s} as $N \rightarrow \infty$.
- The fill distance on the boundary goes to 0 as $M \rightarrow \infty$.
- $\mathcal{H}_K \hookrightarrow H_0^t(\Omega)$ and \mathcal{H}_K is dense in $H_0^t(\Omega)$ (satisfied for Matérn kernels).

Convergence: Main Theorem

Theorem (Convergence to the True Solution)

Let $u^{\gamma, M, N} \in \mathcal{H}_K$ solve the approximate problem, then

$$\lim_{\gamma \rightarrow 0} \lim_{M \rightarrow \infty} \lim_{N \rightarrow \infty} \left\| u^{\gamma, M, N} - u^* \right\|_{H_0^t} = 0.$$




Table of Contents


- 1 Introduction and motivation
- 2 Machine learning for smooth PDEs
- 3 Machine learning for rough PDEs
- 4 Convergence results
- 5 Numerical results**

2D semi-linear PDE


$$\begin{aligned} -\nu \Delta u + u + \sin(\pi u) &= \xi \quad x \in \Omega = (0, 1) \times (0, 1) \\ u &= 0 \quad x \in \partial\Omega \end{aligned}$$

$$u^* \sim \sum_{i,j=1}^{\infty} \frac{u_{ij}}{(i^2 + j^2)^{1+\varepsilon}} 2 \sin(i\pi x) \sin(j\pi y), \quad u_{ij} \sim \mathcal{N}(0, 1) \text{ i.i.d.} \quad \xi \in H^{-1}(\Omega).$$

Solution u



Forcing term ξ



2D semi-linear PDE: pointwise loss

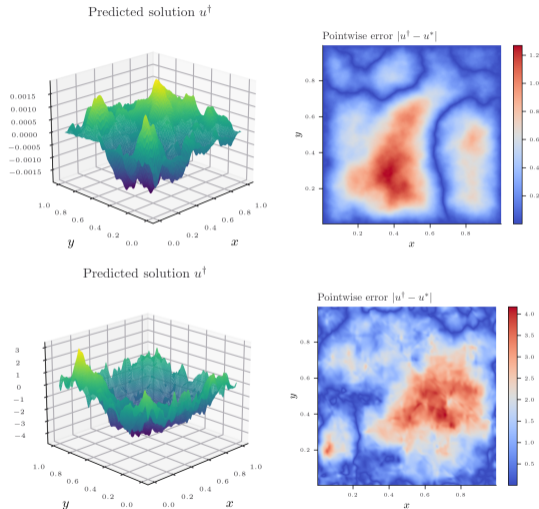


Figure: Kernel method (top), PINN (bottom) ≥ 1.00 relative L^2 error.

2D semi-linear PDE: H^{-1} loss

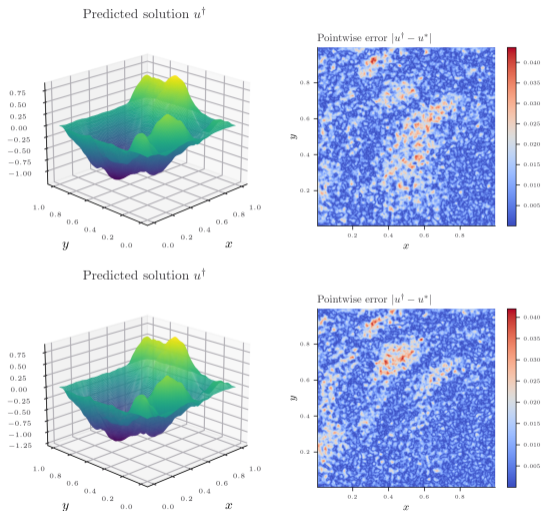


Figure: Kernel method (top), PINN (bottom) ≈ 0.02 relative L^2 error.

Choosing the right norm is very important

When $\xi \in H^{-s}$, $s > 0.95$.

Norm H^{-s}	$s = 0.0$ (L^2)	$s = 0.95$	$s = 0.96$	$s = 1.0$	$s = 2.0$
PINN error	0.312	0.0896	0.0469	0.0469	0.0740

Table: Relative L^2 error for different choices of Sobolev norms as the loss function

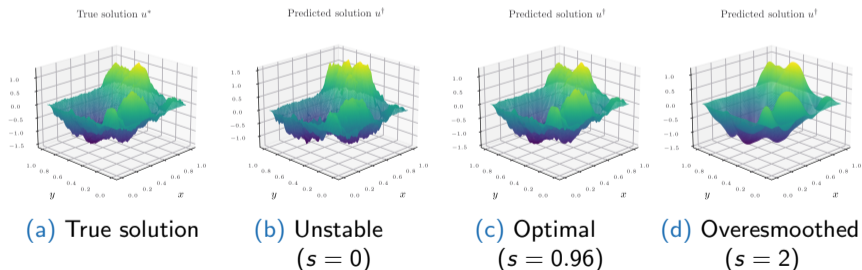


Figure: Effect of different norms on the recovered solution.

Time-dependent: stochastic Allen-Cahn equation

$$\partial_t u = \nu \Delta u + u - u^3 + \sigma \xi \quad \text{in } (0, 1)$$




Figure: Stochastic Allen-Cahn equation.

Conclusion

We propose a kernel-based framework for solving PDEs with irregular forcing terms.

Our **theoretical contributions**:

- We extend machine learning-based solvers to PDEs with weaker norms than L^2 (solving the PDE in weak form).
- We leverage the RKHS structure to provide theoretical guarantees of convergence.
- Applies to linear and non-linear PDEs.

Conclusion

We propose a kernel-based framework for solving PDEs with irregular forcing terms.

Our **theoretical contributions**:

- We extend machine learning-based solvers to PDEs with weaker norms than L^2 (solving the PDE in weak form).
- We leverage the RKHS structure to provide theoretical guarantees of convergence.
- Applies to linear and non-linear PDEs.

Our **computational contributions**:

- We provide an efficient approximation of the negative Sobolev norm.
- We show numerically that this approach is effective for kernel methods and PINNs.
- We provide empirical (and some theoretical) error rates.

R. Baptista, E. Calvello, M. Darcy, H. Owhadi, A. M. Stuart, and X. Yang. *Kernel Methods for Solving Rough PDEs*. 2024.

mdarcy@caltech.edu

References I






-  Batlle, Pau et al. (2023). *Error Analysis of Kernel/GP Methods for Nonlinear and Parametric PDEs*. arXiv: 2305.04962 [math.NA].
-  Chen, Yifan et al. (2021). “Solving and learning nonlinear PDEs with Gaussian processes”. In: *Journal of Computational Physics*.
-  Owhadi, Houman and Clint Scovel (Oct. 2019). *Operator-Adapted Wavelets, Fast Solvers, and Numerical Homogenization*. Cambridge University Press.
-  Raissi, M., P. Perdikaris, and G.E. Karniadakis (2019). “Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations”. In: *Journal of Computational Physics* 378, pp. 686–707.
-  Tancik, Matthew et al. (2020). *Fourier Features Let Networks Learn High Frequency Functions in Low Dimensional Domains*. arXiv: 2006.10739.

Table of Contents

⑥ Appendix: RPDE

⑦ Bonus: numerical method

Bonus: Negative Sobolev Norm approximation

We need to compute an approximation of the negative Sobolev norm:

$$\|f\|_{H^{-s}} \approx \|f\|_{\Phi^N}$$

Bonus: Negative Sobolev Norm approximation

We need to compute an approximation of the negative Sobolev norm:

$$\|f\|_{H^{-s}} \approx \|f\|_{\Phi^N}$$

The recipe:

- 1 Choose a test space $\Phi^N := \text{span}\{\varphi_i\}_{i=1}^N$ to approximate H^s (ex: Haar basis).

Bonus: Negative Sobolev Norm approximation

We need to compute an approximation of the negative Sobolev norm:

$$\|f\|_{H^{-s}} \approx \|f\|_{\Phi^N}$$

The recipe:

- 1 Choose a test space $\Phi^N := \text{span}\{\varphi_i\}_{i=1}^N$ to approximate H^s (ex: Haar basis).
- 2 Measure f against the test space:

$$[f, \varphi] := \left(\int f\varphi_1, \int f\varphi_2, \dots, \int f\varphi_n \right)$$

Bonus: Negative Sobolev Norm approximation

We need to compute an approximation of the negative Sobolev norm:

$$\|f\|_{H^{-s}} \approx \|f\|_{\Phi^N}$$

The recipe:

- 1 Choose a test space $\Phi^N := \text{span}\{\varphi_i\}_{i=1}^N$ to approximate H^s (ex: Haar basis).
- 2 Measure f against the test space:

$$[f, \varphi] := \left(\int f\varphi_1, \int f\varphi_2, \dots, \int f\varphi_n \right)$$

- 3 Compute the stiffness matrix $A \in \mathbb{R}^{N \times N}$, $A_{i,j} := \int_{\Omega} \varphi_i (-\Delta)^s \varphi_j$ and its inverse A^{-1} (can be done efficiently [Owhadi et al.]).

Bonus: Negative Sobolev Norm approximation

We need to compute an approximation of the negative Sobolev norm:

$$\|f\|_{H^{-s}} \approx \|f\|_{\Phi^N}$$

The recipe:

- 1 Choose a test space $\Phi^N := \text{span}\{\varphi_i\}_{i=1}^N$ to approximate H^s (ex: Haar basis).
- 2 Measure f against the test space:

$$[f, \varphi] := \left(\int f \varphi_1, \int f \varphi_2, \dots, \int f \varphi_n \right)$$

- 3 Compute the stiffness matrix $A \in \mathbb{R}^{N \times N}$, $A_{i,j} := \int_{\Omega} \varphi_i (-\Delta)^s \varphi_j$ and its inverse A^{-1} (can be done efficiently [Owhadi et al.]).
- 4 Define

$$\|f\|_{\Phi^N} := \sqrt{[f, \varphi]^T A^{-1} [f, \varphi]}.$$

Finite dimensional problem

Finite dimensional problem

$$u^{\gamma, N, M} = \arg \min_{u \in \mathcal{H}_K} [\mathcal{P}(u) - \xi, \varphi] A^{-1} [\mathcal{P}(u) - \xi, \varphi] + \sum_{j=1}^m |u(x_j)|^2 + \gamma \|u\|_{\mathcal{H}_K}^2$$

Weak solution

This method can be interpreted as solving the PDE in weak form:

$$[\mathcal{P}(u), \varphi_i] = [\xi, \varphi_i] \quad i = 1, \dots, M.$$

This problem can be solved efficiently with the representer theorem and a non-linear least squares optimization techniques such as a variant of the Gauss-Newton algorithm.

Table of Contents

6 Appendix: RPDE


7 Bonus: numerical method

1D Poisson PDE


$$\begin{aligned} -\nu \Delta u + u &= \xi & x \in [0, 1] \\ u &= 0 & x \in \{0, 1\} \end{aligned}$$

$$\xi \sim \sum_{j=1}^{\infty} \xi_j \sqrt{2} \sin(\pi j x), \quad \xi_j \sim \mathcal{N}(0, 1) \text{ i.i.d.}$$


Here $\xi \in H^{-s}(\Omega)$ for any $s > \frac{1}{2}$ and $u^* \in H_0^t(\Omega)$ for $t < \frac{3}{2}$.




(a) Forcing term



(b) Solution



(c) Pointwise error



(d) L^2 error